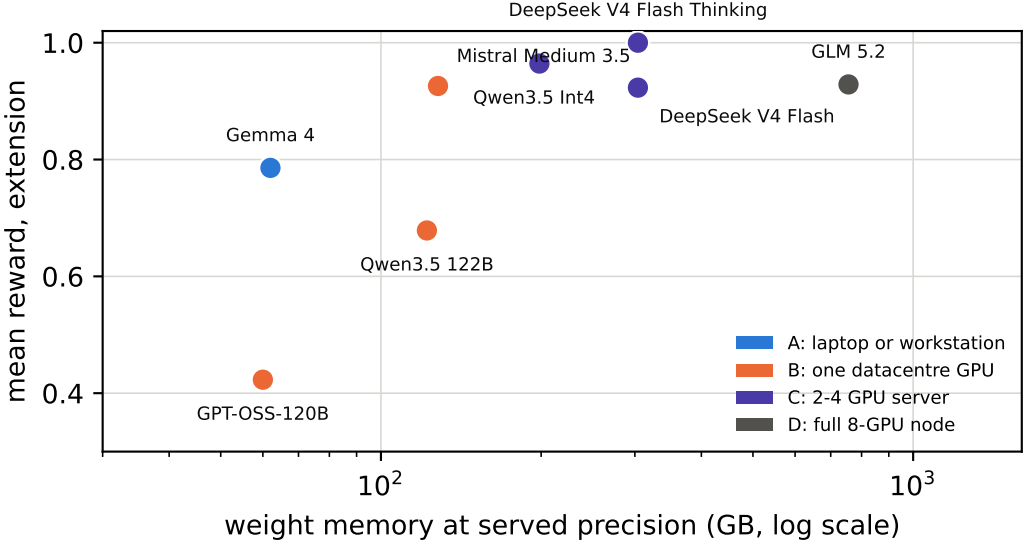
DeepSeek V4 Flash Thinking
Mistral Medium 3.5
GLM 5.2
Qwen3.5 Int4
DeepSeek V4 Flash
Gemma 4
Qwen3.5 122B
GPT-OSS-120B
A: laptop or workstation
B: one datacentre GPU
C: 2-4 GPU server
D: full 8-GPU node
mean reward, extension
1.0
0.8
0.6
0.4
10^2
10^3
weight memory at served precision (GB, log scale)