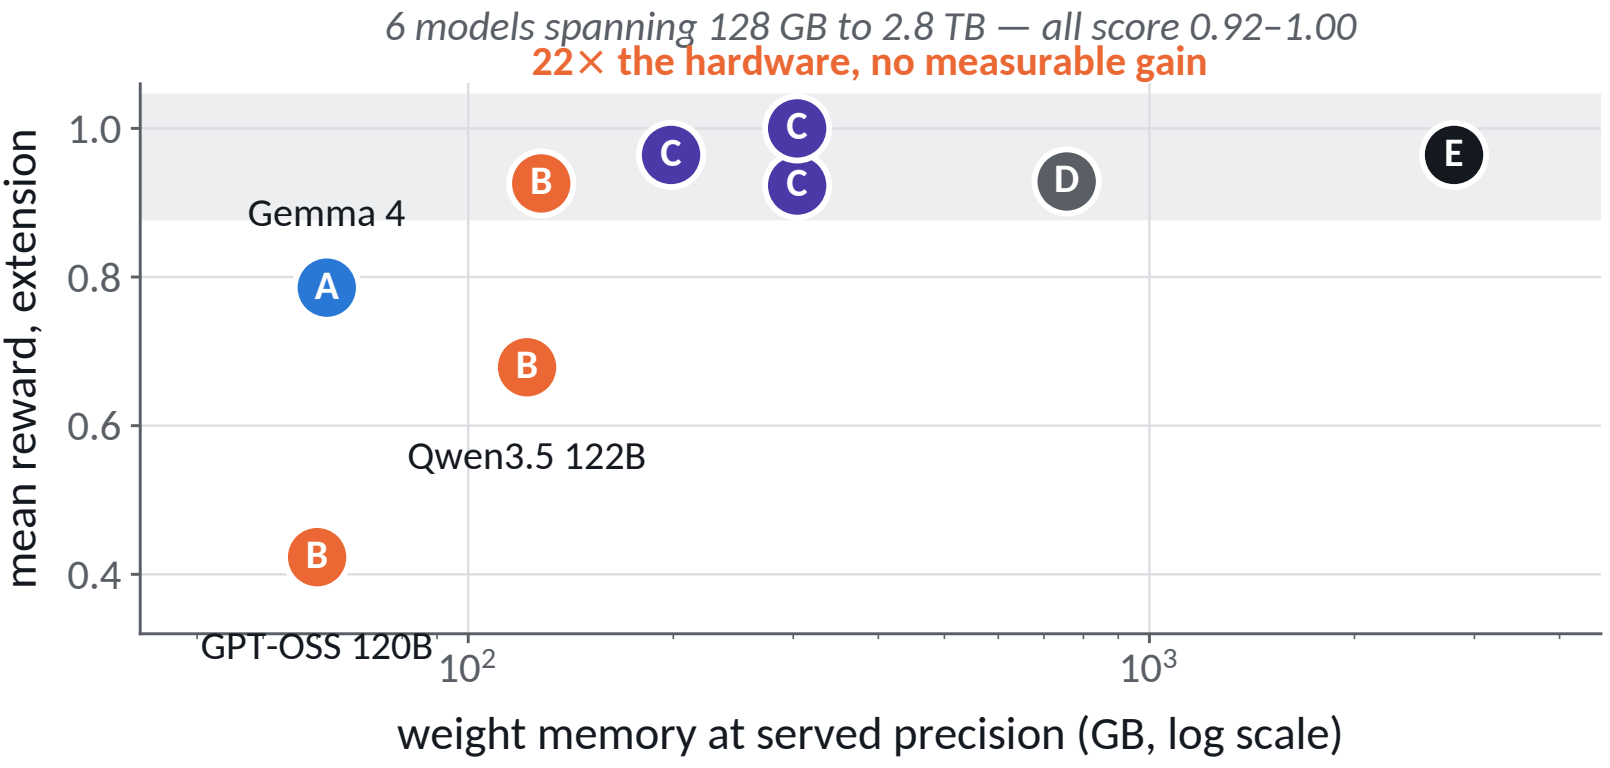
6 models spanning 128 GB to 2.8 TB — all score 0.92–1.00
22× the hardware, no measurable gain
mean reward, extension
1.0
0.8
0.6
0.4
Gemma 4
A
B
C
C
C
D
E
B
Qwen3.5 122B
B
GPT-OSS 120B
10^2
10^3
weight memory at served precision (GB, log scale)